

# Continuous\_Control

December 29, 2020

## 1 Continuous Control

---

You are welcome to use this coding environment to train your agent for the project. Follow the instructions below to get started!

### 1.0.1 1. Start the Environment

Run the next code cell to install a few packages. This line will take a few minutes to run!

```
In [3]: !pip -q install ./python
```

```
# !pip install -U pip
```

```
tensorflow 1.7.1 has requirement numpy>=1.13.3, but you'll have numpy 1.12.1 which is incompatible  
ipython 6.5.0 has requirement prompt-toolkit<2.0.0,>=1.0.15, but you'll have prompt-toolkit 3.0.
```

The environments corresponding to both versions of the environment are already saved in the Workspace and can be accessed at the file paths provided below.

Please select one of the two options below for loading the environment.

```
In [4]: from unityagents import UnityEnvironment  
import numpy as np  
from collections import deque  
import matplotlib.pyplot as plt  
import numpy as np  
import random  
import time  
import torch  
from unityagents import UnityEnvironment
```

```
%matplotlib inline
```

```
# select this option to load version 1 (with a single agent) of the environment
```

```
env = UnityEnvironment(file_name='/data/Reacher_One_Linux_NoVis/Reacher_One_Linux_NoVis.
```

```
# select this option to load version 2 (with 20 agents) of the environment
```

```
# env = UnityEnvironment(file_name='/data/Reacher_Linux_NoVis/Reacher.x86_64')
```

```

INFO:unityagents:
'Academy' started successfully!
Unity Academy name: Academy
    Number of Brains: 1
    Number of External Brains : 1
    Lesson number : 0
    Reset Parameters :
        goal_speed -> 1.0
        goal_size -> 5.0
Unity brain name: ReacherBrain
    Number of Visual Observations (per agent): 0
    Vector Observation space type: continuous
    Vector Observation space size (per agent): 33
    Number of stacked Vector Observation: 1
    Vector Action space type: continuous
    Vector Action space size (per agent): 4
    Vector Action descriptions: , , ,

```

Environments contain *brains* which are responsible for deciding the actions of their associated agents. Here we check for the first brain available, and set it as the default brain we will be controlling from Python.

```

In [9]: # get the default brain
        brain_name = env.brain_names[0]
        brain = env.brains[brain_name]

```

## 1.0.2 2. Examine the State and Action Spaces

Run the code cell below to print some information about the environment.

```

In [10]: # reset the environment
        env_info = env.reset(train_mode=True)[brain_name]

        # number of agents
        num_agents = len(env_info.agents)
        print('Number of agents:', num_agents)

        # size of each action
        action_size = brain.vector_action_space_size
        print('Size of each action:', action_size)

        # examine the state space
        states = env_info.vector_observations
        state_size = states.shape[1]
        print('There are {} agents. Each observes a state with length: {}'.format(states.shape[0], state_size))
        print('The state for the first agent looks like:', states[0])

```

Number of agents: 1  
 Size of each action: 4  
 There are 1 agents. Each observes a state with length: 33  
 The state for the first agent looks like: [ 0.00000000e+00 -4.00000000e+00 0.00000000e+00 1.00000000e+00 -0.00000000e+00 -0.00000000e+00 -4.37113883e-08 0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00 1.00000000e+00 -0.00000000e+00 -0.00000000e+00 -4.37113883e-08 0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00 -6.30408478e+00 -1.00000000e+00 -4.92529202e+00 0.00000000e+00 1.00000000e+00 0.00000000e+00 -5.33014059e-01]

### 1.0.3 3. Take Random Actions in the Environment

In the next code cell, you will learn how to use the Python API to control the agent and receive feedback from the environment.

Note that **in this coding environment, you will not be able to watch the agents while they are training**, and you should set `train_mode=True` to restart the environment.

```
In [8]: env_info = env.reset(train_mode=False)[brain_name]      # reset the environment
        states = env_info.vector_observations                  # get the current state (for each agent)
        scores = np.zeros(num_agents)                         # initialize the score (for each agent)
        while True:
            actions = np.random.randn(num_agents, action_size) # select an action (for each agent)
            actions = np.clip(actions, -1, 1)                  # all actions between -1 and 1
            env_info = env.step(actions)[brain_name]           # send all actions to the environment
            next_states = env_info.vector_observations           # get next state (for each agent)
            rewards = env_info.rewards                          # get reward (for each agent)
            dones = env_info.local_done                         # see if episode finished
            scores += env_info.rewards                          # update the score (for each agent)
            states = next_states                                # roll over states to next time step
            if np.any(dones):                                   # exit loop if episode finished
                break
        print('Total score (averaged over agents) this episode: {}'.format(np.mean(scores)))
```

NameError

Traceback (most recent call last)

```
<ipython-input-8-b273097e0cc3> in <module>()
    1 env_info = env.reset(train_mode=False)[brain_name]      # reset the environment
    2 states = env_info.vector_observations                    # get the current state (for each agent)
----> 3 scores = np.zeros(num_agents)                         # initialize the score (for each agent)
    4 while True:
    5     actions = np.random.randn(num_agents, action_size) # select an action (for each agent)
```

```
NameError: name 'num_agents' is not defined
```

When finished, you can close the environment.

#### 1.0.4 4. It's Your Turn!

Now it's your turn to train your own agent to solve the environment! A few **important notes**: - When training the environment, set `train_mode=True`, so that the line for resetting the environment looks like the following:

```
env_info = env.reset(train_mode=True)[brain_name]
```

- To structure your work, you're welcome to work directly in this Jupyter notebook, or you might like to start over with a new file! You can see the list of files in the workspace by clicking on *Jupyter* in the top left corner of the notebook.
- In this coding environment, you will not be able to watch the agents while they are training. However, *after training the agents*, you can download the saved model weights to watch the agents on your own machine!

```
In [5]: import random
import copy
from collections import deque, namedtuple

import numpy as np
import matplotlib.pyplot as plt

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

from unityagents import UnityEnvironment

DEVICE = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

def hidden_init(layer):
    fan_in = layer.weight.data.size()[0]
    lim = 1. / np.sqrt(fan_in)
    return -lim, lim

class Actor(nn.Module):
    def __init__(self, state_size, action_size, fc1_units, fc2_units):
        super(Actor, self).__init__()
        self.fc1 = nn.Linear(state_size, fc1_units)
```

```

        self.bn1 = nn.BatchNorm1d(fc1_units)
        self.fc2 = nn.Linear(fc1_units, fc2_units)
        self.bn2 = nn.BatchNorm1d(fc2_units)
        self.fc3 = nn.Linear(fc2_units, action_size)
        self.reset_parameters()

    def reset_parameters(self):
        self.fc1.weight.data.uniform_(*hidden_init(self.fc1))
        self.fc2.weight.data.uniform_(*hidden_init(self.fc2))
        self.fc3.weight.data.uniform_(-3e-3, 3e-3)

    def forward(self, state):
        if state.dim() == 1:
            state = torch.unsqueeze(state, 0)

        x = self.bn1(F.relu(self.fc1(state)))
        x = self.bn2(F.relu(self.fc2(x)))
        return F.tanh(self.fc3(x))

class Critic(nn.Module):
    def __init__(self, state_size, action_size, fc1_units, fc2_units):
        super(Critic, self).__init__()
        self.fc1 = nn.Linear(state_size, fc1_units)
        self.bn1 = nn.BatchNorm1d(fc1_units)
        self.fc2 = nn.Linear(fc1_units + action_size, fc2_units)
        self.fc3 = nn.Linear(fc2_units, 1)
        self.reset_parameters()

    def reset_parameters(self):
        self.fc1.weight.data.uniform_(*hidden_init(self.fc1))
        self.fc2.weight.data.uniform_(*hidden_init(self.fc2))
        self.fc3.weight.data.uniform_(-3e-3, 3e-3)

    def forward(self, state, action):
        if state.dim() == 1:
            state = torch.unsqueeze(state, 0)

        x = self.bn1(F.relu(self.fc1(state)))
        x = F.relu(self.fc2(torch.cat((x, action), dim=1)))
        return self.fc3(x)

```

```

Experience = namedtuple('Experience', 'state action reward next_state done')

```

```

class Replay:
    def __init__(self, action_size, buffer_size, batch_size):

```

```

        self.action_size = action_size
        self.buffer = deque(maxlen=buffer_size)
        self.batch_size = batch_size

    def add(self, state, action, reward, next_state, done):
        experience = Experience(state, action, reward, next_state, done)
        self.buffer.append(experience)

    def sample(self):
        experiences = random.sample(self.buffer, k=self.batch_size)

        states = torch.from_numpy(np.vstack([e.state for e in experiences])).float().to(
        actions = torch.from_numpy(np.vstack([e.action for e in experiences])).float().to(
        rewards = torch.from_numpy(np.vstack([e.reward for e in experiences])).float().to(
        next_states = torch.from_numpy(np.vstack([e.next_state for e in experiences])).f
        dones = torch.from_numpy(np.vstack([e.done for e in experiences])).astype(np.uint

        return states, actions, rewards, next_states, dones

    def __len__(self):
        return len(self.buffer)

class Agent:
    def __init__(self, config):
        self.config = config

        self.online_actor = config.actor_fn().to(DEVICE)
        self.target_actor = config.actor_fn().to(DEVICE)
        self.actor_opt = config.actor_opt_fn(self.online_actor.parameters())

        self.online_critic = config.critic_fn().to(DEVICE)
        self.target_critic = config.critic_fn().to(DEVICE)
        self.critic_opt = config.critic_opt_fn(self.online_critic.parameters())

        self.noise = config.noise_fn()
        self.replay = config.replay_fn()

    def step(self, state, action, reward, next_state, done):
        self.replay.add(state, action, reward, next_state, done)

        if len(self.replay) > self.replay.batch_size:
            self.learn()

    def act(self, state, add_noise=True):
        state = torch.from_numpy(state).float().to(DEVICE)

        self.online_actor.eval()

```

```

with torch.no_grad():
    action = self.online_actor(state).cpu().data.numpy()

self.online_actor.train()

if add_noise:
    action += self.noise.sample()

return np.clip(action, -1, 1)

def reset(self):
    self.noise.reset()

def learn(self):
    states, actions, rewards, next_states, dones = self.replay.sample()

    # Update online critic model
    # Predict actions for the next states with the target actor model
    target_next_actions = self.target_actor(next_states)
    # Compute Q values for the next states and actions with the target critic model
    target_next_qs = self.target_critic(next_states, target_next_actions)
    # Compute target Q values for the current states using the Bellman equation
    target_qs = rewards + (self.config.discount * target_next_qs * (1 - dones))
    # Compute Q values for the current states and actions with the online critic model
    online_qs = self.online_critic(states, actions)
    # Compute and minimize the online critic loss
    critic_loss = F.mse_loss(online_qs, target_qs)
    self.critic_opt.zero_grad()
    critic_loss.backward()
    torch.nn.utils.clip_grad_norm_(self.online_critic.parameters(), 1)
    self.critic_opt.step()

    # Update online actor model
    # Predict actions for current states from the online actor model
    online_actions = self.online_actor(states)
    # Compute and minimize the online actor loss
    actor_loss = -self.online_critic(states, online_actions).mean()
    self.actor_opt.zero_grad()
    actor_loss.backward()
    self.actor_opt.step()

    # Update target critic and actor models
    self.soft_update(self.online_critic, self.target_critic)
    self.soft_update(self.online_actor, self.target_actor)

def soft_update(self, online_model, target_model):
    for target_param, online_param in zip(target_model.parameters(), online_model.pa

```

```

        target_param.data.copy_(self.config.target_mix * online_param.data + (1.0 -

class OrnsteinUhlenbeck:
    def __init__(self, size, mu, theta, sigma):
        self.state = None
        self.mu = mu * np.ones(size)
        self.theta = theta
        self.sigma = sigma
        self.reset()

    def reset(self):
        self.state = copy.copy(self.mu)

    def sample(self):
        x = self.state
        dx = self.theta * (self.mu - x) + self.sigma * np.array([random.random() for _ in range(size)])
        self.state = x + dx
        return self.state

def run(agent):
    config = agent.config
    scores_deque = deque(maxlen=100)
    scores = []

    for episode in range(1, config.max_episodes + 1):
        agent.reset()
        score = 0

        env_info = config.env.reset(train_mode=True)[config.brain_name]
        state = env_info.vector_observations[0]

        for step in range(config.max_steps):
            action = agent.act(state)
            env_info = config.env.step(action)[config.brain_name]
            next_state = env_info.vector_observations[0]
            reward = env_info.rewards[0]
            done = env_info.local_done[0]

            agent.step(state, action, reward, next_state, done)

            score += reward
            state = next_state

        if done:
            break

```



```

        scores.append(score)
        scores_deque.append(score)
        mean_score = np.mean(scores_deque)

        print('\rEpisode {} \tAverage Score: {:.2f} \tScore: {:.2f}'.format(episode, mean_score, score))

        if mean_score >= config.goal_score:
            break

    torch.save(agent.online_actor.state_dict(), config.actor_path)
    torch.save(agent.online_critic.state_dict(), config.critic_path)

    fig, ax = plt.subplots()
    ax.plot(np.arange(1, len(scores) + 1), scores)
    ax.set_ylabel('Score')
    ax.set_xlabel('Episode #')
    fig.savefig(config.scores_path)
    plt.show()

class Config:
    def __init__(self, seed):
        self.seed = seed
        random.seed(seed)
        torch.manual_seed(seed)

        self.env = None
        self.brain_name = None
        self.state_size = None
        self.action_size = None
        self.actor_fn = None
        self.actor_opt_fn = None
        self.critic_fn = None
        self.critic_opt_fn = None
        self.replay_fn = None
        self.noise_fn = None
        self.discount = None
        self.target_mix = None

        self.max_episodes = None
        self.max_steps = None

        self.actor_path = None
        self.critic_path = None
        self.scores_path = None

def main():

```

```

config = Config(seed=6)
config.env =UnityEnvironment(file_name='/data/Reacher_One_Linux_NoVis/Reacher_One_Li
config.brain_name = config.env.brain_names[0]
config.state_size = config.env.brains[config.brain_name].vector_observation_space_si
config.action_size = config.env.brains[config.brain_name].vector_action_space_size

config.actor_fn = lambda: Actor(config.state_size, config.action_size, fc1_units=256)
config.actor_opt_fn = lambda params: optim.Adam(params, lr=3e-4)

config.critic_fn = lambda: Critic(config.state_size, config.action_size, fc1_units=2
config.critic_opt_fn = lambda params: optim.Adam(params, lr=3e-4)

config.replay_fn = lambda: Replay(config.action_size, buffer_size=int(1e6), batch_si
config.noise_fn = lambda: OrnsteinUhlenbeck(config.action_size, mu=0., theta=0.15, s

config.discount = 0.99
config.target_mix = 1e-3

config.max_episodes = int(1000)
config.max_steps = int(1e6)
config.goal_score = 30

config.actor_path = 'actor.pth'
config.critic_path = 'critic.pth'
config.scores_path = 'scores.png'

agent = Agent(config)
run(agent)

```

```

if __name__ == '__main__':
    main()

```

Episode 1	Average Score: 1.10	Score: 1.10
Episode 2	Average Score: 0.81	Score: 0.53
Episode 3	Average Score: 0.81	Score: 0.79
Episode 4	Average Score: 0.82	Score: 0.85
Episode 5	Average Score: 0.94	Score: 1.41
Episode 6	Average Score: 0.89	Score: 0.69
Episode 7	Average Score: 0.77	Score: 0.00
Episode 8	Average Score: 0.99	Score: 2.58
Episode 9	Average Score: 1.03	Score: 1.29
Episode 10	Average Score: 0.97	Score: 0.46
Episode 11	Average Score: 0.98	Score: 1.04
Episode 12	Average Score: 0.98	Score: 0.97
Episode 13	Average Score: 1.02	Score: 1.56
Episode 14	Average Score: 1.13	Score: 2.60
Episode 15	Average Score: 1.09	Score: 0.48

Episode 16	Average Score: 1.14	Score: 1.91
Episode 17	Average Score: 1.17	Score: 1.58
Episode 18	Average Score: 1.12	Score: 0.37
Episode 19	Average Score: 1.12	Score: 1.08
Episode 20	Average Score: 1.28	Score: 4.38
Episode 21	Average Score: 1.31	Score: 1.77
Episode 22	Average Score: 1.31	Score: 1.33
Episode 23	Average Score: 1.37	Score: 2.71
Episode 24	Average Score: 1.38	Score: 1.74
Episode 25	Average Score: 1.43	Score: 2.42
Episode 26	Average Score: 1.47	Score: 2.55
Episode 27	Average Score: 1.53	Score: 3.14
Episode 28	Average Score: 1.66	Score: 5.24
Episode 29	Average Score: 1.73	Score: 3.55
Episode 30	Average Score: 1.82	Score: 4.40
Episode 31	Average Score: 1.88	Score: 3.80
Episode 32	Average Score: 1.90	Score: 2.58
Episode 33	Average Score: 1.95	Score: 3.46
Episode 34	Average Score: 2.02	Score: 4.43
Episode 35	Average Score: 2.11	Score: 5.01
Episode 36	Average Score: 2.11	Score: 1.99
Episode 37	Average Score: 2.26	Score: 7.69
Episode 38	Average Score: 2.28	Score: 3.14
Episode 39	Average Score: 2.33	Score: 4.43
Episode 40	Average Score: 2.36	Score: 3.55
Episode 41	Average Score: 2.36	Score: 2.16
Episode 42	Average Score: 2.50	Score: 8.16
Episode 43	Average Score: 2.56	Score: 5.21
Episode 44	Average Score: 2.55	Score: 2.15
Episode 45	Average Score: 2.66	Score: 7.33
Episode 46	Average Score: 2.68	Score: 3.78
Episode 47	Average Score: 2.72	Score: 4.28
Episode 48	Average Score: 2.74	Score: 3.69
Episode 49	Average Score: 2.72	Score: 1.80
Episode 50	Average Score: 2.83	Score: 8.25
Episode 51	Average Score: 2.85	Score: 4.19
Episode 52	Average Score: 3.02	Score: 11.20
Episode 53	Average Score: 3.03	Score: 3.74
Episode 54	Average Score: 3.04	Score: 3.80
Episode 55	Average Score: 3.10	Score: 5.99
Episode 56	Average Score: 3.13	Score: 4.87
Episode 57	Average Score: 3.23	Score: 9.09
Episode 58	Average Score: 3.31	Score: 7.57
Episode 59	Average Score: 3.49	Score: 14.32
Episode 60	Average Score: 3.54	Score: 6.51
Episode 61	Average Score: 3.63	Score: 8.73
Episode 62	Average Score: 3.81	Score: 14.90
Episode 63	Average Score: 3.85	Score: 6.50

Episode 64	Average Score: 3.90	Score: 6.76
Episode 65	Average Score: 4.01	Score: 10.90
Episode 66	Average Score: 4.19	Score: 15.98
Episode 67	Average Score: 4.38	Score: 17.09
Episode 68	Average Score: 4.49	Score: 11.47
Episode 69	Average Score: 4.51	Score: 5.94
Episode 70	Average Score: 4.66	Score: 15.48
Episode 71	Average Score: 4.68	Score: 5.87
Episode 72	Average Score: 4.80	Score: 13.62
Episode 73	Average Score: 4.81	Score: 5.54
Episode 74	Average Score: 4.87	Score: 8.55
Episode 75	Average Score: 4.90	Score: 7.51
Episode 76	Average Score: 5.13	Score: 22.69
Episode 77	Average Score: 5.18	Score: 8.91
Episode 78	Average Score: 5.17	Score: 3.84
Episode 79	Average Score: 5.18	Score: 6.24
Episode 80	Average Score: 5.16	Score: 3.72
Episode 81	Average Score: 5.38	Score: 22.58
Episode 82	Average Score: 5.49	Score: 14.34
Episode 83	Average Score: 5.60	Score: 14.76
Episode 84	Average Score: 5.76	Score: 19.33
Episode 85	Average Score: 5.91	Score: 18.45
Episode 86	Average Score: 5.95	Score: 9.48
Episode 87	Average Score: 6.03	Score: 12.63
Episode 88	Average Score: 6.15	Score: 16.47
Episode 89	Average Score: 6.19	Score: 10.05
Episode 90	Average Score: 6.21	Score: 8.20
Episode 91	Average Score: 6.48	Score: 30.29
Episode 92	Average Score: 6.65	Score: 22.25
Episode 93	Average Score: 6.74	Score: 15.35
Episode 94	Average Score: 6.93	Score: 24.71
Episode 95	Average Score: 7.08	Score: 20.65
Episode 96	Average Score: 7.16	Score: 14.64
Episode 97	Average Score: 7.19	Score: 10.20
Episode 98	Average Score: 7.27	Score: 15.43
Episode 99	Average Score: 7.28	Score: 8.43
Episode 100	Average Score: 7.42	Score: 21.24
Episode 101	Average Score: 7.54	Score: 12.63
Episode 102	Average Score: 7.70	Score: 16.27
Episode 103	Average Score: 7.82	Score: 12.71
Episode 104	Average Score: 7.98	Score: 16.89
Episode 105	Average Score: 8.08	Score: 11.84
Episode 106	Average Score: 8.14	Score: 6.25
Episode 107	Average Score: 8.31	Score: 17.43
Episode 108	Average Score: 8.46	Score: 17.99
Episode 109	Average Score: 8.50	Score: 4.89
Episode 110	Average Score: 8.76	Score: 26.80
Episode 111	Average Score: 8.91	Score: 15.57

Episode 112	Average Score: 9.12	Score: 21.67
Episode 113	Average Score: 9.39	Score: 28.53
Episode 114	Average Score: 9.64	Score: 27.77
Episode 115	Average Score: 9.74	Score: 10.43
Episode 116	Average Score: 9.98	Score: 26.33
Episode 117	Average Score: 10.25	Score: 28.92
Episode 118	Average Score: 10.59	Score: 33.73
Episode 119	Average Score: 10.83	Score: 25.10
Episode 120	Average Score: 10.88	Score: 9.29
Episode 121	Average Score: 11.11	Score: 25.15
Episode 122	Average Score: 11.31	Score: 21.54
Episode 123	Average Score: 11.56	Score: 27.56
Episode 124	Average Score: 11.74	Score: 19.45
Episode 125	Average Score: 11.97	Score: 25.13
Episode 126	Average Score: 12.10	Score: 15.89
Episode 127	Average Score: 12.36	Score: 29.58
Episode 128	Average Score: 12.43	Score: 11.92
Episode 129	Average Score: 12.75	Score: 35.09
Episode 130	Average Score: 13.09	Score: 38.42
Episode 131	Average Score: 13.43	Score: 38.15
Episode 132	Average Score: 13.62	Score: 21.60
Episode 133	Average Score: 13.96	Score: 37.84
Episode 134	Average Score: 14.24	Score: 31.60
Episode 135	Average Score: 14.58	Score: 39.32
Episode 136	Average Score: 14.67	Score: 10.78
Episode 137	Average Score: 14.93	Score: 33.79
Episode 138	Average Score: 15.19	Score: 29.43
Episode 139	Average Score: 15.37	Score: 22.40
Episode 140	Average Score: 15.72	Score: 38.39
Episode 141	Average Score: 15.84	Score: 14.02
Episode 142	Average Score: 16.06	Score: 30.33
Episode 143	Average Score: 16.26	Score: 25.37
Episode 144	Average Score: 16.56	Score: 32.30
Episode 145	Average Score: 16.76	Score: 26.91
Episode 146	Average Score: 17.11	Score: 38.87
Episode 147	Average Score: 17.33	Score: 25.92
Episode 148	Average Score: 17.66	Score: 37.55
Episode 149	Average Score: 17.79	Score: 13.98
Episode 150	Average Score: 18.07	Score: 36.24
Episode 151	Average Score: 18.37	Score: 34.81
Episode 152	Average Score: 18.52	Score: 25.78
Episode 153	Average Score: 18.72	Score: 23.93
Episode 154	Average Score: 18.93	Score: 24.57
Episode 155	Average Score: 19.10	Score: 23.47
Episode 156	Average Score: 19.42	Score: 36.52
Episode 157	Average Score: 19.63	Score: 30.01
Episode 158	Average Score: 19.91	Score: 35.42
Episode 159	Average Score: 20.06	Score: 29.98

Episode 160	Average Score: 20.34	Score: 33.99
Episode 161	Average Score: 20.62	Score: 37.41
Episode 162	Average Score: 20.81	Score: 33.62
Episode 163	Average Score: 21.09	Score: 34.25
Episode 164	Average Score: 21.33	Score: 31.33
Episode 165	Average Score: 21.61	Score: 38.79
Episode 166	Average Score: 21.73	Score: 27.82
Episode 167	Average Score: 21.95	Score: 38.94
Episode 168	Average Score: 22.10	Score: 26.79
Episode 169	Average Score: 22.42	Score: 37.41
Episode 170	Average Score: 22.57	Score: 30.83
Episode 171	Average Score: 22.88	Score: 36.97
Episode 172	Average Score: 23.06	Score: 31.34
Episode 173	Average Score: 23.33	Score: 32.06
Episode 174	Average Score: 23.54	Score: 30.39
Episode 175	Average Score: 23.82	Score: 34.68
Episode 176	Average Score: 23.83	Score: 24.25
Episode 177	Average Score: 24.13	Score: 38.40
Episode 178	Average Score: 24.45	Score: 35.88
Episode 179	Average Score: 24.72	Score: 33.79
Episode 180	Average Score: 25.08	Score: 39.16
Episode 181	Average Score: 25.18	Score: 32.94
Episode 182	Average Score: 25.39	Score: 35.48
Episode 183	Average Score: 25.63	Score: 38.15
Episode 184	Average Score: 25.81	Score: 38.01
Episode 185	Average Score: 26.02	Score: 39.50
Episode 186	Average Score: 26.30	Score: 37.53
Episode 187	Average Score: 26.57	Score: 38.89
Episode 188	Average Score: 26.79	Score: 38.85
Episode 189	Average Score: 27.03	Score: 34.33
Episode 190	Average Score: 27.33	Score: 38.47
Episode 191	Average Score: 27.38	Score: 35.09
Episode 192	Average Score: 27.55	Score: 39.33
Episode 193	Average Score: 27.76	Score: 36.42
Episode 194	Average Score: 27.91	Score: 39.09
Episode 195	Average Score: 28.03	Score: 33.06
Episode 196	Average Score: 28.26	Score: 37.34
Episode 197	Average Score: 28.41	Score: 25.51
Episode 198	Average Score: 28.60	Score: 33.87
Episode 199	Average Score: 28.86	Score: 34.46
Episode 200	Average Score: 28.99	Score: 34.94
Episode 201	Average Score: 29.14	Score: 27.29
Episode 202	Average Score: 29.32	Score: 34.69
Episode 203	Average Score: 29.52	Score: 32.13
Episode 204	Average Score: 29.66	Score: 30.54
Episode 205	Average Score: 29.92	Score: 38.15
Episode 206	Average Score: 30.25	Score: 39.49

